

Exercises 2 Online Learning - R Code

Natalia Zuniga-Garcia

9/27/2017

Part C

First, I use the functions coded for Exercise 1 Part 2. These functions estimate the success probability (sigmoid), the negative log-likelihood, and the gradient of the negative log-likelihood.

```
# Function 1. Estimate the Success Probability
wi.estimate <- function(b, X){
  # Inputs:
  # b regression parameter P x 1
  # X Matrix of features N x P
  # Output:
  # wi Success Probability N x 1

  w <- 1/(1+exp(-X %*% b))
  w = pmax(w, 1E-6); w = pmin(w, 1-1E-6);
  return(w)
}

# Function 2. Estimate the negative log likelihood
Negll <- function(b, y, X, m){
  # Input:
  # b regression parameter P x 1
  # y vector of response N x 1
  # X Matrix of features N x P
  # m number of trial for the ith case
  # Output:
  # Negative log likelihood of binomial distribution

  ll <- -sum(dbinom(y, m, wi.estimate(b, X), log = TRUE))
  return(ll)
}

# Function 3. Estimate the gradient of the negative log likelihood
Gradll <- function(b, y, X, m){
  # Input:
  # b regression parameter P x 1
  # y vector of response N x 1
  # X Matrix of features N x P
  # m number of trial for the ith case
  # Output:
  # gradient of l(b) P x 1

  grad.ll <- as.numeric(y - m * wi.estimate(b, X)) * X
  return(-colSums(grad.ll))
}
```

Now, I code up the Stochastic Gradient Descent.

```

# Function 4. Stochastic Gradient Descend
SGradDes <- function(y, X, beta, m, iter, epsilon, alpha){
  # Input:
  # y vector of response N x 1
  # X Matrix of features N x P
  # beta: initial guess of beta P x 1
  # m: number of trial for the ith case
  # iter: maximum iterations allowed if it doesn't converge
  # epsilon: minimum error allowed for convergency criteria
  # alpha: step size (gamma based on class document)
  # Output:
  # Negative log likelihood per iteration
  # Using stochastic gradient descend
  # b regression parameter P x 1

  # Initial Iteration (GUESS)
  betas = array(NA, dim=c(iter, ncol(X)))
  betas[1,] = beta
  ll = array(NA, dim = iter)
  ll[1] = Negll(betas[1,], y, X, m)

  # Iterations
  for (i in 2:iter){
    r <- sample(nrow(X), 1) # Draw a random sample with replacement
    grad <- Gradll(betas[i-1,], matrix(y[r], nrow=1), matrix(X[r,], nrow=1), m)

    # Gradient Descend
    betas[i,] <- betas[i-1,] - alpha * grad
    ll[i] <- Negll(betas[i,], y, X, m)
    # Note that I am tracking the full objective function in every step
    # Thus I am not touching only one point per iteration as SGD intended
    # I'm doing this with the objective of develop a learning process for myself

    # Checking for Convergence
    error = abs((ll[i] - ll[i-1])/(ll[i-1] + epsilon))
    if (error < epsilon){
      cat('Stochastic Gradient Descend has converged in iterations:', i)
      ll <- ll[1:i]
      betas <- betas[1:i,]
      break;
    } else if (i == iter && error >= epsilon){
      print('Stochastic Gradient Descend has not converged')
      break;
    }
  }
  return(list("Negll" = ll, "beta" = betas[i,]))
}

```

Now I evaluate different alpha (step size) values using a simulation with wdbc database.

```

a <- c(0.001, 0.01, 0.025, 0.1, 0.25, 1)
for (j in 1:length(a)){
  sgd <- SGradDes(y, X, beta, m, iter, epsilon, alpha=a[j])
  llSGD <- as.matrix(unlist(sgd[1], use.names = FALSE))
}

```


```

betaSGD <-  as.matrix(unlist(sgd[2], use.names = FALSE))

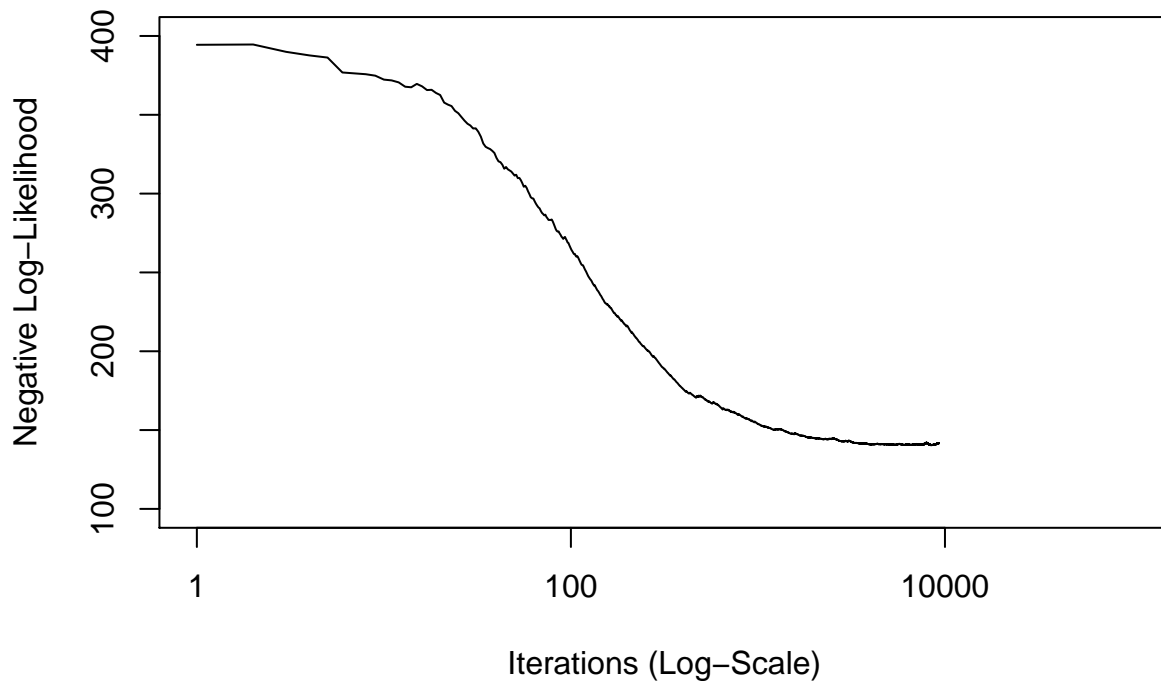
#png(filename=paste('F1PC',a[j], '.png'),width=15,height=12,units="cm",res=200)
plot(1:length(llSGD),llSGD,type="l",col="black", xlab="Iterations (Log-Scale)", ylab="Negative Log-Likelihood",
#dev.off()
}

```

Stochastic Gradient Descend has converged in iterations: 59718




Stochastic Gradient Descend has converged in iterations: 9278



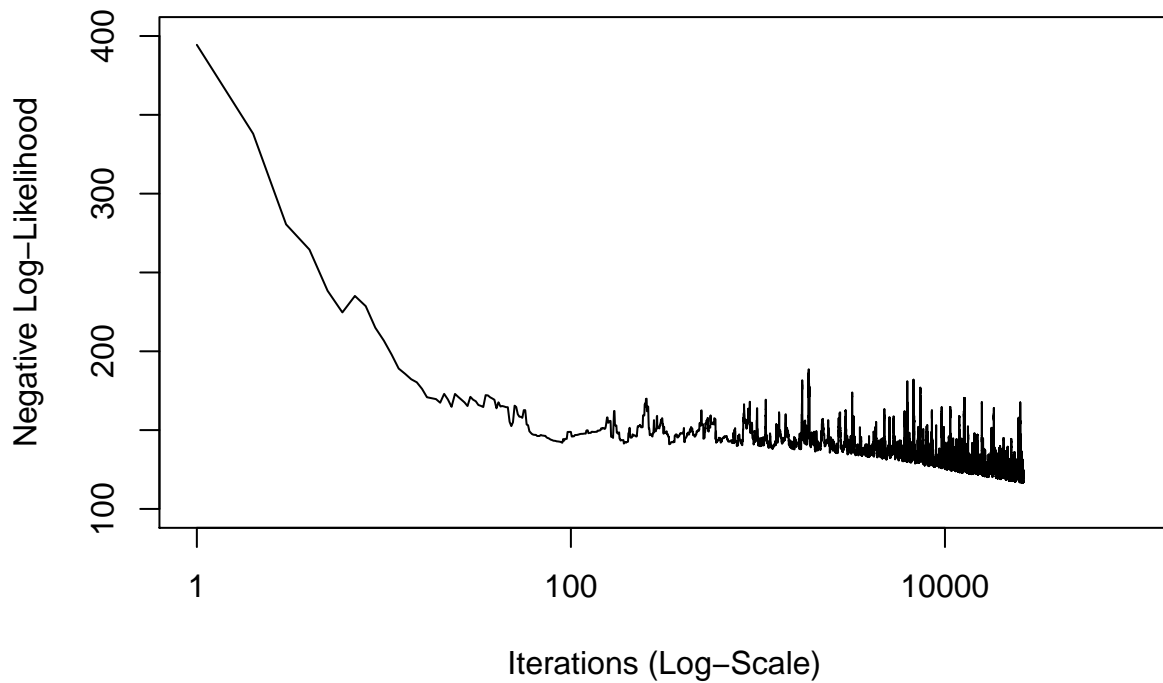
```
## Stochastic Gradient Descend has converged in iterations: 9318
```




```
## Stochastic Gradient Descend has converged in iterations: 10930
```



```
## Stochastic Gradient Descend has converged in iterations: 26329
```



Stochastic Gradient Descent has converged in iterations: 88220



Part D

In this part I implement SGD using a decaying step based on Robbins-Monro rule.

```
# Function 5. Stochastic Gradient Descent Decaying Step Size
SGradDesRM <- function(y, X, beta, m, iter, epsilon, C, t, alpha){
  # Input:
```

```

# y vector of response N x 1
# X Matrix of features N x P
# beta: initial guess of beta P x 1
# m: number of trial for the ith case
# iter: maximum iterations allowed if it doesn't converge
# epsilon: minimum error allowed for convergency criteria
# C: constant in the Robbins-Monro rule
# t: t0 prior number of steps Robbins-Monro rule
# alpha: learning rate Robbins-Monro rule
# Output:
# Negative log likelihood per iteration
# Using stochastic gradient descend
# with decaying step size
# b regression parameter P x 1

# Initial Iteration (GUESS)
betas = array(NA, dim=c(iter, ncol(X)))
betas[1,] = beta
ll = array(NA, dim = iter)
ll[1] = Negll(betas[1,], y, X, m)

# Iterations
for (i in 2:iter){
  r <- sample(nrow(X), 1) # Draw a random sample with replacement
  grad <- Gradll(betas[i-1,], matrix(y[r], nrow=1), matrix(X[r,], nrow=1), m)

  # Gradient Descend
  step <- C*(i+t)^(-alpha) # Robbins-Monro rule
  betas[i,] <- betas[i-1,] - step * grad
  ll[i] <- Negll(betas[i,], y, X, m)

  # Checking for Convergence
  error = abs((ll[i] - ll[i-1])/(ll[i-1] + epsilon))
  if (error < epsilon){
    cat('Stochastic Gradient Descend has converged in iterations:',i)
    ll <- ll[1:i]
    betas <- betas[1:i,]
    break;
  } else if (i == iter && error >= epsilon){
    print('Stochastic Gradient Descend has not converged')
    break;
  }
}
return(list("Negll" = ll, "beta" = betas[i,]))
}

```

Now I evaluate different C values.

```

t <- 1
alpha <- 0.6
C <- c(0.01,0.1,0.4,0.6,0.8,1)
for (j in 1:length(a)){
  sgdrm = SGradDesRM(y, X, beta, m, iter, epsilon,C[j], t, alpha)
  llSGDrm = as.matrix(unlist(sgdrm[1], use.names = FALSE))
}

```


```

betaSGDrM = as.matrix(unlist(sgdrm[2], use.names = FALSE))


#png(filename=paste('F2PD',C[j], '.png'),width=15,height=12,units="cm",res=200)
plot(1:length(llSGDrM),llSGDrM,type="l",col="black", xlab="Iterations (Log-Scale)", ylab="Negative Log-Likelihood",
#dev.off()
}

```


Stochastic Gradient Descend has converged in iterations: 26531



Stochastic Gradient Descend has converged in iterations: 57520



Stochastic Gradient Descend has converged in iterations: 22451




Stochastic Gradient Descend has converged in iterations: 14933



Stochastic Gradient Descend has converged in iterations: 11887



Stochastic Gradient Descend has converged in iterations: 9422



Finally I evaluate different alpha values.


```
C <- 0.1
alpha <- c(0.5,0.6,0.7,0.8,0.9,1)
for (j in 1:length(alpha)){
  sgdrm = SGradDesRM(y, X, beta, m, iter, epsilon,C, t, alpha[j])
  llSGDrM = as.matrix(unlist(sgdrm[1], use.names = FALSE))
  betaSGDrM = as.matrix(unlist(sgdrm[2], use.names = FALSE))
}
```

```


#png(filename=paste('F3PD',alpha[j],'.png'),width=15,height=12,units="cm",res=200)
plot(1:length(llSGDrm),llSGDrm,type="l",col="black", xlab="Iterations (Log-Scale)", ylab="Negative Log-Likelihood",
#dev.off()
}

```

Stochastic Gradient Descend has converged in iterations: 48745




Stochastic Gradient Descend has converged in iterations: 58550




[1] "Stochastic Gradient Descend has not converged"



```
## [1] "Stochastic Gradient Descend has not converged"
```



```
## Stochastic Gradient Descend has converged in iterations: 64267
```



Stochastic Gradient Descent has converged in iterations: 64288

